UNIVERSITÄT PADERBORN
Die Universität der Informationsgesellschaft

Faculty for Computer Science, Electrical Engineering and Mathematics
Department of Computer Science
Research Group Data Science

Master's Thesis

Submitted to the Data Science Research Group
in Partial Fullfilment of the Requirements for the Degree of

Master of Science

Grammar-based Compression of RDF Graphs

by
PHILIP FRERK

Thesis Supervisors:

Prof. Dr. Axel-Cyrille Ngonga Ngomo

Prof. Dr. Stefan Böttcher

Paderborn, May 2019

Erklärung

Ich versichere, dass ich die Arbeit ohne fremde Hilfe und ohne Benutzung anderer als der angegebenen Quellen angefertigt habe und dass die Arbeit in gleicher oder ähnlicher Form noch keiner anderen Prüfungsbehörde vorgelegen hat und von dieser als Teil einer Prüfungsleistung angenommen worden ist. Alle Ausführungen, die wörtlich oder sinngemäß übernommen worden sind, sind als solche gekennzeichnet.

_____   _____
          Ort, Datum   Unterschrift

Abstract. We present a full documentation of the Paderborn University Computer Science thesis template (UPB-CS-TT) and how to use it. This document also serves as a demonstrator to show what documents UPB-CS-TT produces. Have fun!

Contents

1

Introduction

1.1 Motivation

The majority of data on the Internet is unstructured because it is mainly text. That makes it difficult for machines to extract knowledge from the data and answer specific queries. In the context of the Semantic Web, an attempt is made to build a knowledge base using structured data. Here the data is stored as graphs. A graph is an often used data structure which consists of nodes and edges. The edges connect the different nodes. This thesis will focus on knowledge graphs, which are about expressing knowledge or facts. A possible format for the such graphs is the Resource Description Framework (RDF)[1], in which a graph is represented by triples of the form ($subject$, $predicate$, $object$), where $subject$ and $object$ are nodes and $predicate$ is an edge of the graph. A triple is typically used to express a certain fact.

In reality, such knowledge graphs can become very large with millions or even billions of nodes and edges. The following three use cases can then become very hard or even infeasible: storage, transmission and processing. One way to circumvent this problem is to use compression.

The currently most popular compressor for RDF data is HDT [FMPG+13]. Here the data is made smaller by a compact representation of the triples.

There is a completely different compression technique which is called grammar-based compression. This method has so far been tested very little for RDF graphs. As the name suggests, such a compression is based on the principle of a formal grammar, with productions that can be nested among each other. There are not yet many grammar-based compression algorithms for graphs. One of the best known is GraphRePair [MP18]. This is a compressor that works for any type of graph and is therefore also applicable for RDF. Why the algorithm is particularly suitable for RDF will be explained later on.

In this thesis it shall be investigated to what extent grammar-based compression is applicable for RDF and whether it delivers even better results than HDT.

1.2 General Idea

As already mentioned, there are essentially three use cases for RDF data:

1. Storage

---

[1]https://www.w3.org/RDF/

2. Transmission

3. Processing/Consumption

The idea is that you can make all these use cases easier / possible by compression, especially if one is dealing with very large amounts of data. In the first two use cases, compression can be helpful by reducing the size of the data. Thus the data can be stored more compactly and above all can be transferred faster, which occurs frequently in reality and can become a problem due to possibly slow transfer speeds . here possibly concrete numbers
The third use case (Processing/Consumption) is about reading or writing access to data. The more common case of read access is typically the execution of queries on RDF data. For example, you want to find out which other nodes a given node is connected to. But there are other examples which will be explained in the course of the thesis. A compression algorithm can help in this respect by making it possible to answer such queries directly on the compressed graphs. This may even be faster than on the original data due to the reduced size.
The aim of the thesis is to compare the two compressors (HDT and GRP) with regard to these use cases and finally determine which approach works better for which use cases.

1.3 Thesis Structure

In chapter 2 the necessary fundamentals are presented. It will mainly be about the two compressors HDT and GRP. In chapter **??** some definitions are introduced, which are needed for the further thesis.
Chapter 3 deals with the first main part of the thesis, the comparison of HDT and GRP.
Then the second main part is presented in chapter 4. It deals with how to further improve the compression ratio by taking advantage of special RDF features.
Chapter 3 and 4 each consist of a theoretical part and an evaluating part, which is supposed to prove the theory.
Finally, the results of the thesis are summarized in chapter 5 and the future work on this topic is presented.

2

Related Work

This chapter gives an introduction to the domain RDF and introduces the different compression algorithms discussed in this thesis.

2.1 RDF

RDF is a format for structured data on the web. It can be formally described as follows. Let the following sets be infinite and mutually disjoint:

- U (URI references, typically represent unique entities or properties)

- B (Blank nodes, represent an arbitrary value, necessary for displaying more complex logical statements)

- L (Literals, represent fixed values, e.g. numbers or strings)

An RDF triple $(s, p, o) \in (U \cup B) \times U \times (U \cup B \cup L)$ displays the statement that the subject s is related to the object o via the predicate p. So a subject can anything but a literal. A predicate can only be a URI and an object can be anything.

2.2 HDT

HDT ([FMPG$^+$13]) is the best known approach to compressing RDF data because it is directly tailored to RDF. The idea behind it is to make RDF's extensive representation more compact, thus reducing memory size. In the normal RDF format (e.g. turtle), each triple is written down after another. This does not only require much space, but is also not good for query performance, since each single triple has to be traversed within a query execution.

It is easy to see that a more compact representation can be achieved by cleverly grouping the triples. That is exactly what HDT does. Imagine for example some triples that all have the same subject s. These would then be written in turtle format as follows:

$$(s, p_1, o_{11}), ..., (s, p_1, o_{1_{n1}}), (s, p_2, o_{2_1}), ...,$$
$$(s, p_2, o_{2_{n2}}), ..., (s, p_k, o_{kn_k})$$

In HDT the triples are then grouped by s.

$$s \rightarrow [(p_1, (o_{1_1}, ..., o_{1_{n_1}})), (p_2, (o_{2_1}, ..., o_{2_{n_2}})), ...,$$
$$(p_k, (o_{k_{n_k}}))]$$

Since all triples have the same subject s, it does not have to be written down at the beginning of each triple anymore. On the second level the triples are then grouped according to the predicates. So all triples with the predicate p_1 come first, then all with the predicate p_2 and so on. Finally only the different objects have to be enumerated since the subject and predicate are already implicitly given. You will now see how HDT implements this mechanism.

First, all URIs and literals (which are usually quite long) are mapped to unique IDs (integers). From now on these IDs will be used. You can see the triples with the IDs in Fig. 2.1 on the left ('Plain Triples').

The above mentioned grouped representation is called 'Compact Triples' and can be seen in Fig. 2.1 in the middle. In the array 'Predicates' you can see the IDs of the predicates. A '0' means that from now on a new subject comes. For example, the first two entries in the 'Predicates' are linked to the first subject, since there is a '0' in the third position in the array.

This works analogously for the objects, in the array 'Objects' the IDs of the objects are listed, and a '0' here means that from then on a new predicate comes. For instance, the first entry in 'Objects' is a '6', so the first derived triple would be $(1, 2, 6)$. The second entry of 'Objects' is a '0' since there is a change from predicate '2' to '3' in 'Predicates'.

The last representation of triples is called 'Bitmap Triples'. This is a slight adaption of 'Compact Triples'. The array S_p has the same content as 'Predicates', but without the zeros. That job is done by the bit-array B_p in which a '1' at position i means that at position i in S_p there comes a change of the subject (this change was previously denoted by a '0' in 'Predicates').

That works analogously for the objects where S_o has the same content as 'Objects', but without zeros.

The advantage of 'Bitmap Triples' is that queries can be executed faster on this representation. Also, the compressed size is slightly smaller than with 'Compact Triples'. Therefore, 'Bitmap Triples' are always used in the current version of HDT.

HDT has also procedures for answering queries, but those will not be mentioned here, because they are not relevant for the thesis. [FMPG$^+$13]
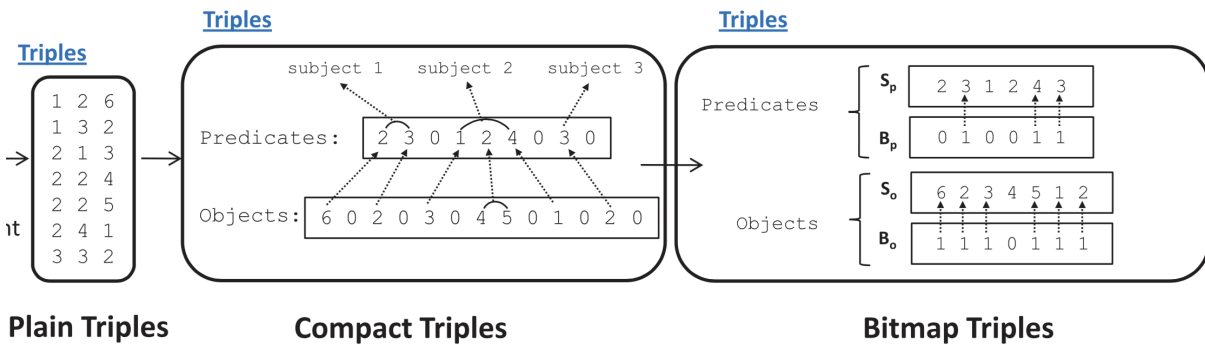

Figure 2.1: Three different representations of triples in HDT

2.3 Grammar-based Graph Compression

In this section we will discuss two different approaches to grammar-based graph compression. First, [Dü16] is introduced and briefly explained why the approach is less suitable for this thesis. Then [MP18], which the thesis will focus on, is introduced.

2.3.1 Dürksen's Algorithm

An approach to grammar-based compression of graphs was developed in [Dü16]. Here the authors assume a hyper graph with node and edge labels. Such a graph can be seen in Fig. 2.2 on the left. To simplify the compression, a so-called transformation is now executed, in which a new node is inserted for each edge e, which then has the same label as e. The original structure of the graph is obtained by connecting two nodes, which were previously connected by an edge, indirectly by the new node. The result of the transformation is a graph, which only has node labels, but no edge labels. Moreover, hyper edges (edges with more than two incident nodes) are no longer present due to the transformation.
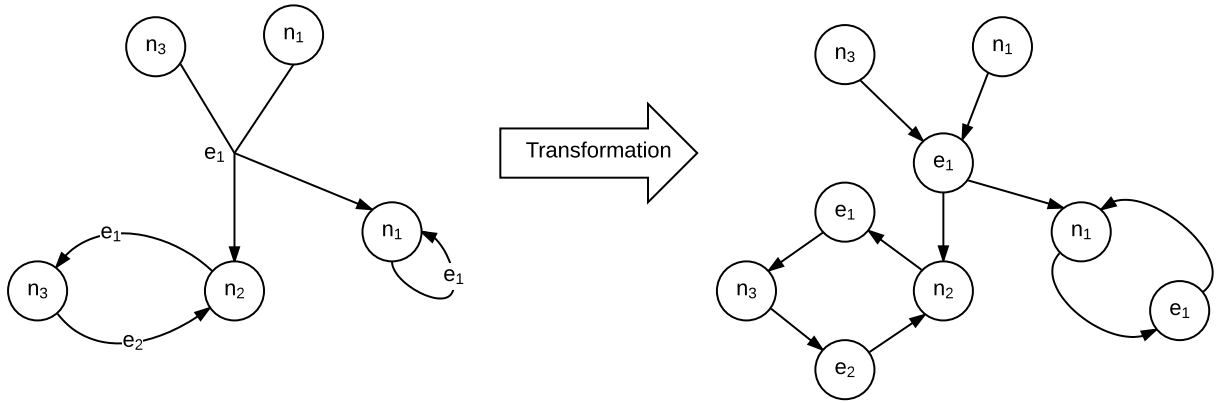

Figure 2.2: Transformation of a labeled graph (each edge is replaced by a new node having the label of the replaced edge).

Thus patterns like in Fig. 2.3 can be replaced. Here it happens twice that a node with the label X is connected with a node with the label Y. This pattern is then stored in a central location and can be referenced via the label X'. Thus only two nodes remain in the compressed graph (with X' as label) and the graph was reduced to a smaller size. There are some details which make this pattern replacement possible, but they are neglected at this point.

Dürksen's algorithm does not seem to be suitable for RDF, because it is based on the fact that there are several nodes with the same label. However, since in RDF a node represents an entity that normally does not occur twice, Dürksen's basic assumption is not fulfilled in RDF.

In addition, Dürksen's algorithm is not yet in a mature state, i.e. there is only a rudimentary implementation that does not work reliably for large graphs. In addition, work is currently underway to find a compact representation of such a compressed graph in order to save the data with little memory. [Dü16]

For these reasons we will focus on the compressor GraphRePair, which is presented in chapter 2.3.2.

2.3.2 GraphRePair

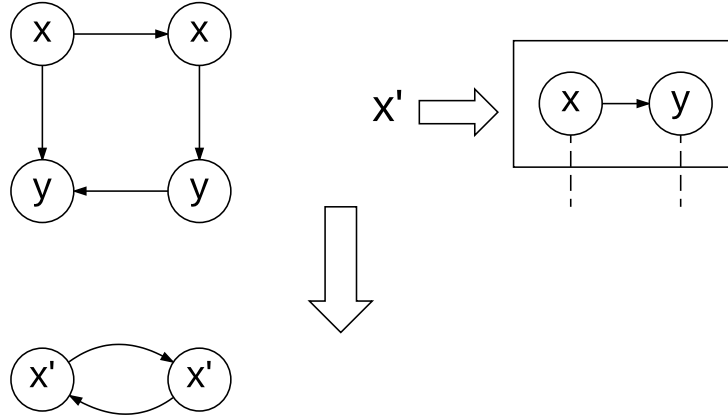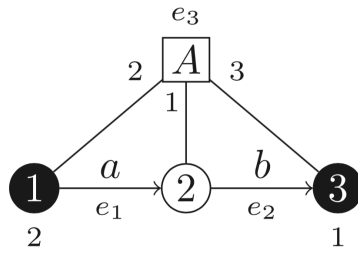This chapter deals with GraphRePair, the compressor from [MP18].

Figure 2.3: Replacement of two occurrences of the pattern $X \to Y$ by nodes with the label X'.

Figure 2.4

Foundations

For a set M, $M^+ = \{x_1 \cdot x_2 \cdot ... \cdot x_n | x_1, x_2, .., x_n \in M\}$ is defined as the set of all non-empty strings of M where $\cdot$ stands for the concatenation of two symbols. $M^* = M^+ \cup \{\epsilon\}$ is similar to M^+, but it also includes the empty string ϵ

A hypergraph over Σ is a tuple $g = (V, E, att, lab, ext)$ where $V = \{1, ..., n\}$ is the set of nodes. $E \subseteq \{(i,j) | i, j \in V\}$ is the set of edges. $att : E \to V^+$ is the attachment mapping, $lab : E \to \Sigma$ is the edge label mapping, and $ext \in V^*$ is a series of external nodes.

A hypergraph does not contain multi-edges, which means for two edges $e_1 \neq e_2$ it holds $att(e_1) \neq att(e_2) \vee lab(e_1) \neq lab(e_2)$.

For a hypergraph $g = (V, E, att, lab, ext)$ $V_g, E_g, att_g, lab_g, ext_g$ are used to refer to its components.

An example for a hypergraph is illustrated in Fig. 2.4. Formally the graph can be described as $V = \{1, 2, 3\}$, $E = \{e_1, e_2, e_3\}$, $att = \{e_1 \mapsto 1 \cdot 2, e_2 \mapsto 2 \cdot 3, e_3 \mapsto 2 \cdot 1 \cdot 3\}$, $lab = \{e_1 \mapsto a, e_2 \mapsto b, e_3 \mapsto A\}$ and $ext = 3 \cdot 1$.

External nodes are black and the numbers below them indicate indices for their position in ext. Analogously, hyper edges (e_3 in this case) have indices indicating the order of the attached nodes. [MP18]

Digram

Digrams are like patterns, the same digram can occur multiple times in a graph. Based on that, GRP reduces a graph's size.

A digram d is a hyper graph with $E_d = \{e_1, e_2\}$ so that the following conditions hold:

1. $\forall v \in V_d : v \in att_d(e_1) \vee v \in att_d(e_2)$

6

2. $\exists v \in V_d : v \in att_d(e_1) \land v \in att_d(e_2)$

3. $ext_d \neq \epsilon$

Condition 1 ensures that all nodes in d are incident to one of the two edges of d. Conditions 2 is used to make sure that there is one 'middle node' incident to both edges of d. Finally, condition 3 ensures that there are external nodes. [MP18]

Digram Occurrence

Let g be a hyper graph and d be a digram with the two edges e_1^d, e_2^d. Let $o = \{e_1, e_2\} \subseteq E_g$ and let V_o be the set of nodes incident with edges in o. Then o is an occurrence of d in g if there exits a bijection $b : V_o \to V_d$ so that for $i \in \{1, 2\}$ and $v \in V_o$ all following conditions hold:

1. $b(v) \in att_d(e_i^d)$ iff $v \in att_g(e_i)$

2. $lab_d(e_i^d) = lab_g(e_i)$

3. $b(v) \in ext_d$ iff $v \in att_g(e)$ for some $e \in E_g \setminus o$

Condition 1 and 2 ensure that the two edges of o form a graph isomorphic to d. Condition 3 makes sure that every external node of d is mapped to a node in g that is incident to at least one edge in g that is not contained in o. [MP18]

Algorithm

Algorithm 1 is the main routine of GraphRePair. The algorithm take a graph as input and returns a grammar whereas N is the set of non terminals, P is the set of productions and $S \in P$ is the start production. It maintains a list of digram occurrences in g. As long as this list contains multiple occurrences for one digram the loop will continue. In the loop, the most frequent digram is found and then replaced and the grammar is extended. After a digram replacement the occurrence list has to be updated, because the graph has now changed and former digram occurrences may not exist anymore. Moreover, new digram occurrences can be present now. This occurrence update is a complex process and will not be discussed here, since it is not relevant for the thesis. [MP18]

Algorithm 1 GraphRePair (Graph $g = (V, E, att, lab, ext)$)

1: $N, P \leftarrow \emptyset$
2: $S \leftarrow g$
3: $L(d) \leftarrow$ list of non-overlapping occurrences of every digram d appearing in g
4: **while** $|L(d)| > 1$ for at least one digram d **do**
5:     $mfd \leftarrow$ most frequent digram
6:     $A \leftarrow$ new non terminal for mfd
7:     Replace every occurrence of mfd in g
8:     $N \leftarrow N \cup \{A\}$
9:     $P \leftarrow P \cup \{A \to mfd\}$
10:     Update the occurrence list L
11: **return** Grammar $G = (N, P, S)$

	0	1	2	3	4	5	6	
0	0	1	1	0	0	0	0	0
1	0	0	0	1	0	0	0	0
2	0	0	0	0	1	0	0	0
3	0	0	0	0	0	1	0	0
4	0	0	0	0	0	0	0	0
5	1	0	0	0	0	0	0	0
6	0	1	1	0	0	0	0	0
	0	0	0	0	0	0	0	0

(a) Adjacency matrix of a graph.

(b) Conceptual K^2-tree with $k = 2$.

$T_1 = 1110$

$T_2 = 1100\ 0010\ 1011$

$L = 0100\ 1001\ 1001\ 0010\ 0100\ 1000$

$L_c = 01\ 11\ 11\ 00\ 01\ 10$
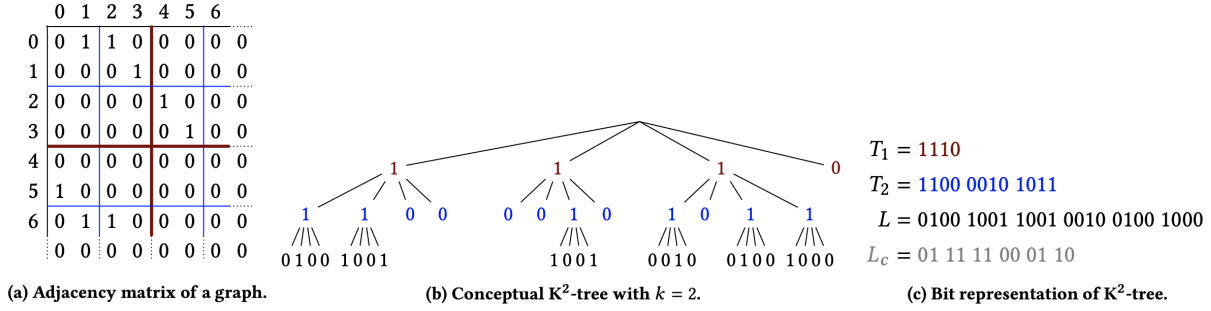
(c) Bit representation of K^2-tree.

Figure 2.5: An adjacency matrix, its k^2 tree representation and the tree's bit representation.

Grammar Encoding

After GRP has constructed a grammar for some graph, this grammar has to be stored in an efficient way. The start production (the graph) is most often much bigger than the other productions and is therefore encoded differently.

The start production is encoded using k^2 trees which is illustrated in Fig. 2.5.

First the matrix is extends with zeros to the next power of two. Then it is partitioned into k^2 equally large partitions ($k = 2$ here).

The tree's root now represents the whole matrix and its child order correspond to the order of the just created partitions. If a partition only contains zeros, a zero leaf is added to the tree (e.g. in Fig. 2.5 (a), partition 4, right bottom). Otherwise a one-node is added and the corresponding condition will be partitioned itself. The recursive procedures continues until there is zero-leaf for each path of the tree.

Afterwards, the tree is represented by bit-strings.

Since the graph has also edge labels, an adjacency matrix and its corresponding tree will be created for each of those labels.

The remaining productions of the grammar are encodes differently, because they are usually quite small compared to the start production. Here so-called δ-codes with variable length are used. Essentially this is a way of displaying numbers as a bit-string in an efficient way (similar to a Huffman Code). To realize that, the right hand side of the components of the production have to represented by numbers, e.g. the edges labels, the node IDs etc. [MP18]

3

GraphRePair vs HDT

This chapter deals with the comparison between HDT and GRP. The main aim is to determine which of the two compressors achieves a lower compression ratio.

First, in Ch. 3.1 some necessary definitions are introduced.

In Ch. 3.2 some basic concepts of the compressors will be discussed again and hypotheses will be made how the compression ratio behaves in relation to a certain structuring of RDF data.

In Ch. 3.3 it will be considered how data can be generated that is structured in the way stated in Ch. 3.2.

Finally in Ch. 3.4 an empirical evaluation of the compression ratios for HDT and GRP for that synthetic RDF data will be done. The evaluation will show that the hypotheses stated in Ch. 3.2 are correct.

3.1 Definitions

This chapter includes some basic definitions that will be used later on.

3.1.1 Compression Ratio

One of the key metrics for a compressor is its compression ratio. The compression ratio depends on the input data and is defined as follows

$$\text{compression ratio} = \frac{\text{compressed size}}{\text{original size}} \in (0, 1]$$

The compression ratio is typically in the $(0, 1]$ interval, since the compressed data will not be larger than the original data (there are cases where this happens, but it does not happen in the compressors considered here). Obviously, the compressed data cannot have a size of 0 or less. Normally the compression ratio is measured at the file size level (in byte). If a different measure is used, it will be mentioned at that point.

3.1.2 (De-)Compression Time

Another key metric of a compressor is its (de-)compression time. This metric also depends on the input data and indicates the run time needed for compression and decompression of the data, respectively. The run time is typically measured in milliseconds.

3.2 Theory

The obvious question is now whether there are certain properties/features that an RDF graph can have, and which have a positive or negative impact on the compression ratio of one or both algorithms.

3.2.1 Relation Between Structure of Data and Compression Ratio

First these features are considered for HDT. Fig. 3.1 is shown again. There you can see that the size of the data becomes smaller if there are only a few subjects. This is the case because the bit-array B_p contains a 1 every time a new subject is considered. For example, if there is only one subject, then B_p consists only of zeros.
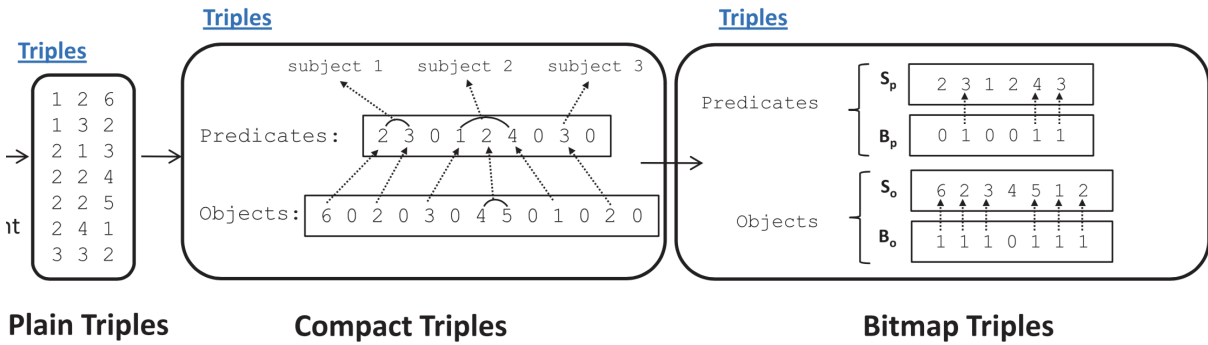

Figure 3.1: Three different representations of triples in HDT

For GRP that feature analysis is more complex. Since GRP constructs grammar rules by using the graph's structure, it is more dependent on the data's structure than HDT. There can be many features that can lead to different constructible grammar rules. Those will not be discussed here. But a simple insight is that GRP's compression ratio will be bigger when there are more different predicates in the graph. This is true, because GRP's grammar rules are based on repeating edge labels. That fact will be considered in the evaluation.

3.2.2 Dictionary Size

As already explained in Ch. 2, HDT divides an RDF file into its header, dictionary and triples component. This is partially also true for GRP except that GRP does not create a header. But it also assigns an ID to every URI or literal and then only works with these IDs. Unfortunately the authors of [MP18] did not work on efficient storing of the dictionary. In GRP the size for the dictionary component is just ignored. Therefore we have to add this size in order to compare GRP with HDT in a fair way. To achieve that we just add the same size to the compressed size of GRP that HDT would need to store the dictionary.

3.3 Synthetic Data Creation

As has just explained, HDT can compress a graph that is similar to a hub pattern (few subjects, many objects) very well. So it gets worse the further away the graph is from this pattern. This corresponds to the authority pattern, where there are few objects but many subjects.
The task now is to create a series of RDF graphs that first correspond to the hub pattern and then continue to change in the direction of the authority pattern. This is illustrated in Fig. 3.2. In this scenario all graphs (G_1 to G_m) have the same size, i.e. the same number of nodes

and edges. G_1 has only one subject connected to all objects. G_2 then has two subjects more and correspondingly 2 objects less. This goes on and on until there is only one object that is connected to all subjects (G_m). The edges are randomly distributed among the nodes, so that all nodes have a similar degree and each node has at least a degree of one.
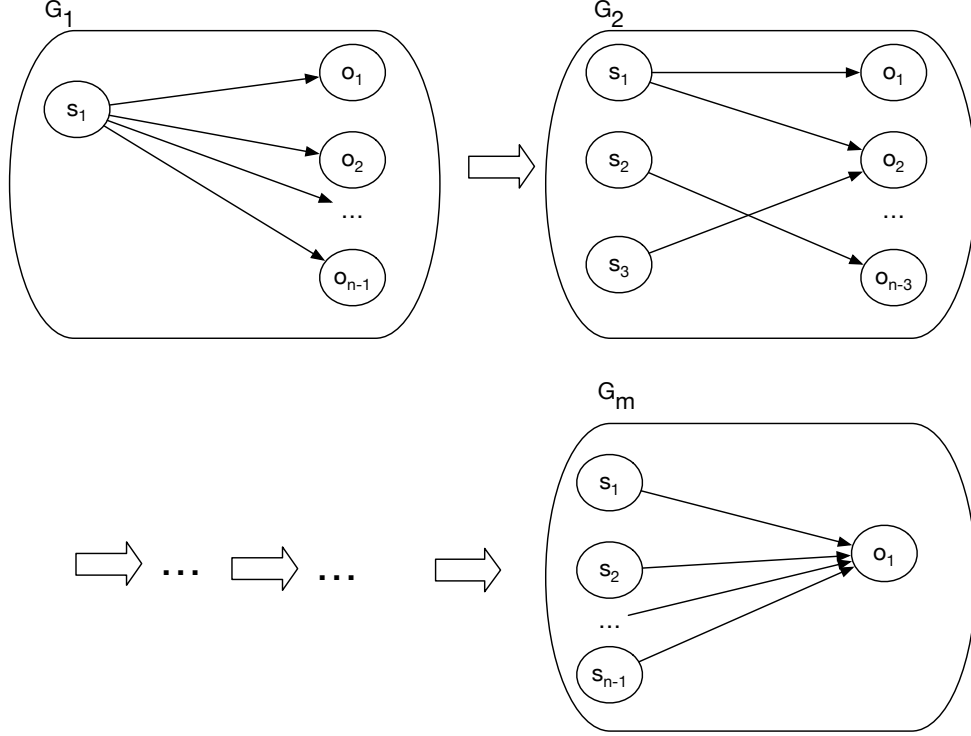

Figure 3.2: Step-by-step transition from hub pattern to authority pattern. The number of nodes n is the same for each graph. The number of edges is also the same for each graph.

It is also ensured that each of the generated files has exactly the same size. This is made possible by ensuring that each URI has the same length. Since the RDF graph also has the same number of triples, the files are of the same size. Since the evaluation compares the compression ratios for the different RDF files, it is important that all files are of the same size to ensure a fair comparison.

A section of such a file (for $G1$) is shown in Fig. 3.3. In that example, there is only one distinct predicate for all triples. The number of predicates is always one at first. The amount of predicates has a similar effect on both compressors and is therefore omitted at first. But in Ch. 3.4 that effect will be discussed in more detail.

Apart from that, blank nodes and literals are not used here. For both compressors, blank nodes and literals are being handled analogously to URI nodes. Therefore they are not needed at this point, in order to show the behavior of HDT and GRP in the above described scenario.

3.4   Evaluation

For the following evaluation HDT-Java 2.0 [1] (the currently newest version) has been used. For GRP the implementation mentioned in [MP18] has been used (written in Scala). It is not open source and has been given to us by the authors of [MP18].

---

[1] https://github.com/rdfhdt/hdt-java/releases/tag/v2.0

```
<http://subject/0000000000> <http://predicate/0000000000> <http://object/0000001036> .
<http://subject/0000000000> <http://predicate/0000000000> <http://object/0000000854> .
<http://subject/0000000000> <http://predicate/0000000000> <http://object/0000000991> .
<http://subject/0000000000> <http://predicate/0000000000> <http://object/0000000450> .
<http://subject/0000000000> <http://predicate/0000000000> <http://object/0000000456> .
<http://subject/0000000000> <http://predicate/0000000000> <http://object/0000000689> .
<http://subject/0000000000> <http://predicate/0000000000> <http://object/0000001166> .
<http://subject/0000000000> <http://predicate/0000000000> <http://object/0000000960> .
```

Figure 3.3: Excerpt from the generated RDF file for G_1 (see Fig. 3.2). Each triple has the same length.

Fig. 3.4a shows the compression ratio for HDT (without dictionary size). As expected, the ratio gets higher the more similar the graph is to the authority pattern. In general it can be said that this effect is quite small. There is only a distance of 0.002 between the minimum and the maximum. This small effect can also be seen by looking at Fig. 3.4b. It can be seen that the size of the dictionary has a much bigger effect, since the compression ratio is now much larger and the curve behavior from Fig. 3.4a is no longer recognizable. It is noticeable that the dictionary size gets bigger when the graph is further away from the star pattern. The dictionary implementation of HDT seems to be more inefficient when there are about as many subject as objects.

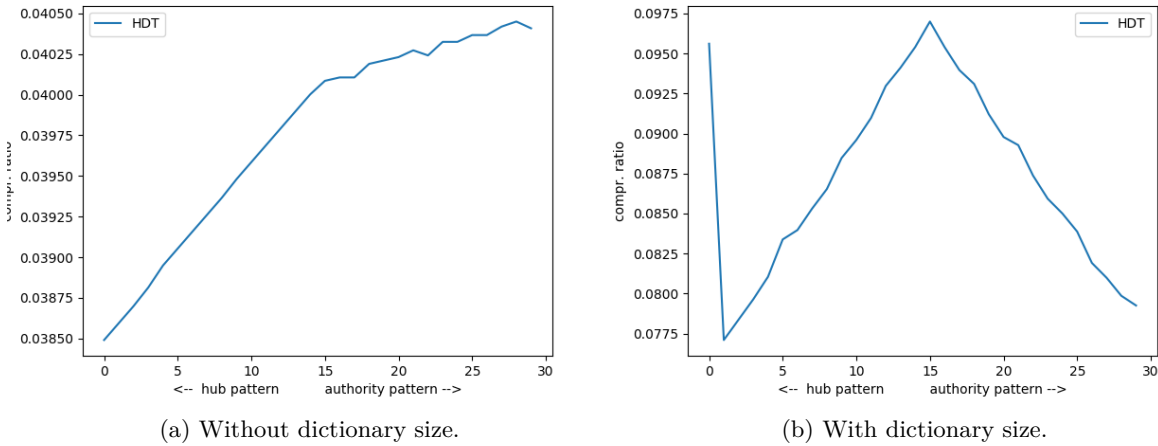(a) Without dictionary size.

(b) With dictionary size.

Figure 3.4: The compression ratios for HDT without and with dictionary sizes.

Next the compression ratio of GRP is considered, which is presented in Fig. 3.5a (without dictionary sizes). Here one can see that GRP has a better compression ratio if the graph is more similar to the star pattern (hub order authority pattern). This property of GRP has also been mentioned in [MP18]. It can also be seen that the effect on the compression ratio is bigger for GRP than for HDT (standard deviation is twice as high for GRP as for HDT). A grammar-based compression is therefore more dependent on the structure of the input data.

When Fig. 3.5b is considered, it can be seen that this curve behaves almost exactly like the one from Fig. 3.4b, since the size of the dictionary accounts for most of the compressed data size.

Finally, Fig. 3.6a and Fig.3.6b show the compression ratios for both algorithms. Since both use the same method to compress the dictionary, the curves in Fig. 3.6b are very similar. However, it becomes clear that GRP compresses better than HDT. In Fig. 3.6a, the ratio of HDT is 31 times higher on average. Of course, this factor becomes much smaller in Fig. 3.6b because the dictionary accounts for most of the memory size. Here the compression ratio of HDT is on
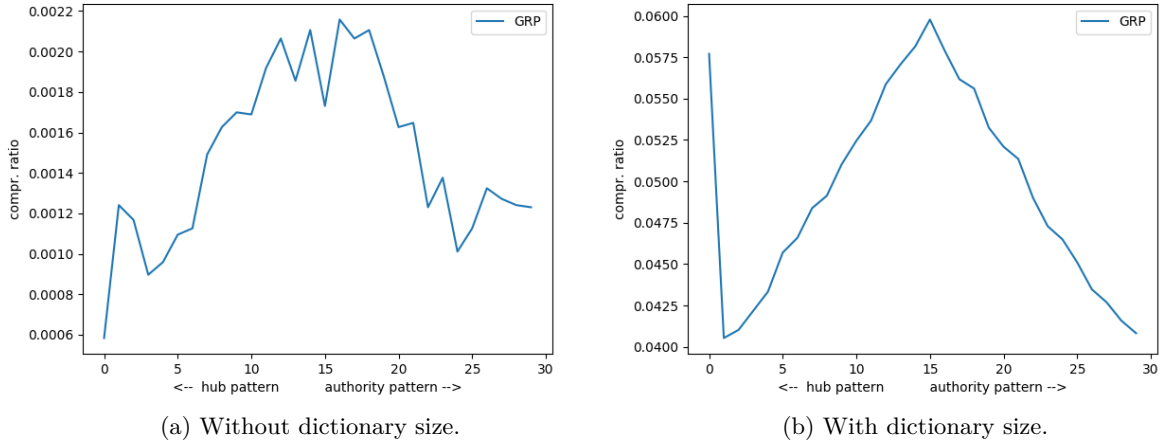
(a) Without dictionary size.

(b) With dictionary size.

Figure 3.5: The compression ratios for GRP without and with dictionary sizes.

average 1.8 times as high as that of GRP.

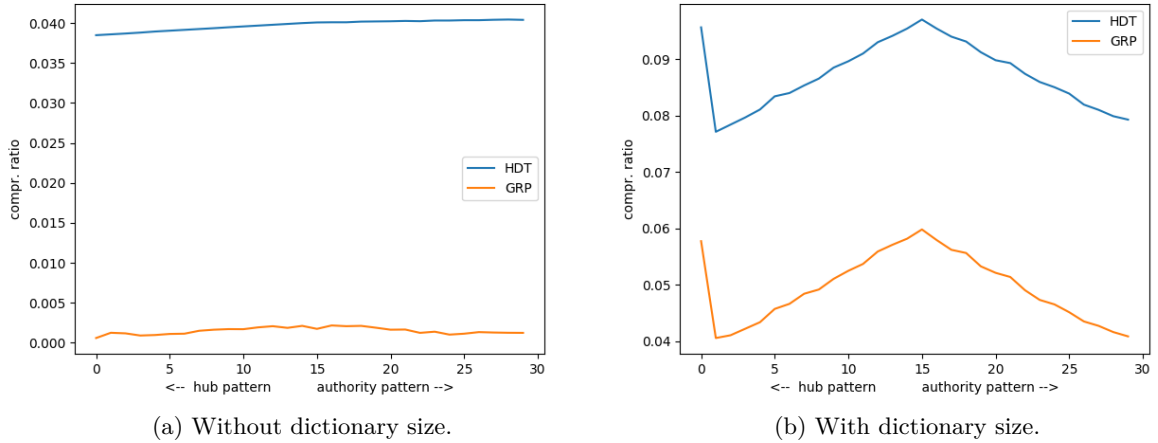(a) Without dictionary size.

(b) With dictionary size.

Figure 3.6: The compression ratios for GRP and HDT without and with dictionary sizes.

One could now argue that only one distinct predicate was used in that scenario and this is beneficial for GRP, as it gets worse as the number of predicates increases. Therefore a further evaluation is made in Fig. 3.7, where 1000 distinct predicates have been used. That is a quite high number considering the number of triples (1199) compared to real RDF data . One can see that the compression ratios are now higher for both compressors, but still GRP's ratio is always smaller than HDT's. HDT's ratio is still 1.7 times higher on average. So, the increasing number of predicates has a similar effect on both algorithms.

Apart from the compression ratio, the run time is also important for the overall performance. Fig. 3.8 shows the average run times of the two algorithms. For this the same scenario with the star pattern (and only one distinct predicate) was used. It has been executed 100 times to get a sophisticated run time measurement, because the run time also depends on the current CPU workload of the computer.

It can be seen that the runtime of GRP is significantly higher than that of GRP. It is on average
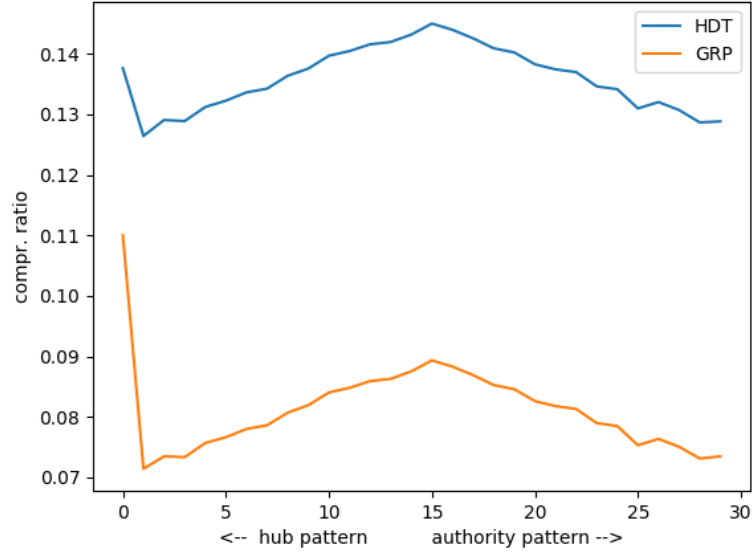
Figure 3.7: The compression ratios for GRP and HDT with dictionary sizes. Graphs have now 1000 distinct predicates.
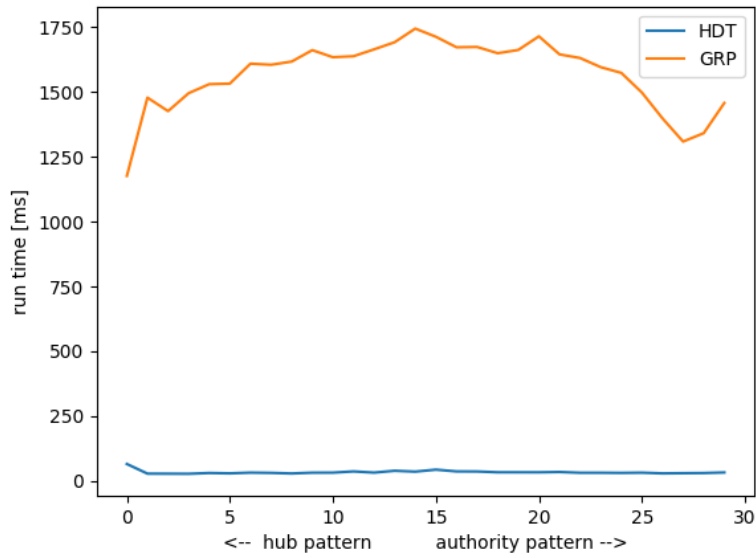

Figure 3.8: Run times of both algorithms (average run time of 100 consecutive executions).

ca. 48 times as high.

However, it should also be noted that the implementation of GRP is rather rudimentary (according to the authors of [MP18]), while that of HDT has been under development for some time. So they are not comparable in terms of quality. Unfortunately, one cannot say at this point whether a more professional implementation of GRP will also be slower than HDT.

In addition one can notice that GRP's run time fluctuates more than that of HDT. GRP has a standard deviation of about 134, while HDT only has a standard deviation of about 7. One reason for this is that GRP, in contrast to HDT, is non-deterministic because of the partly random search order of the graph. On the other hand, the high deviations are also a confirmation of the above mentioned hypothesis that the behavior of GRP depends more on the structure of the input data than HDT does.

4

Using RDF-specific Features

GRP is a compressor made for all kinds of graphs. This also means that it is not specially adapted for RDF. It is therefore possible to use some special features of RDF to achieve a better compression ratio. This is what chapter 4.1 will deal with.

On the other hand, chapter 3 already revealed that most of the compressed size is determined by the dictionary, since saving the dictionary is more complex than saving the graph structure. Chapter 4.2 will therefore focus on how to compress the dictionary even more.

4.1   Ontology Knowledge

4.1.1   Theory

4.1.2   Evaluation

4.2   Dictionary Features

4.2.1   Theory

4.2.2   Evaluation

5

Summary and Discussion

discussion

Bibliography

[Dü16]   Matthias Dürksen. Grammatik-basierte Graphkompression. Bachelor-arbeit, Universität Paderborn, 2016.

[FMPG+13] Javier D. Fernández, Miguel A. Martínez-Prieto, Claudio Gutiérrez, Axel Polleres, and Mario Arias. Binary rdf representation for publication and exchange (hdt). *Web Semantics: Science, Services and Agents on the World Wide Web*, 19:22 – 41, 2013.

[MP18]   Sebastian Maneth and Fabian Peternek. Grammar-based graph compression. *Inf. Syst.*, 76:19–45, 2018.